

Cancer/Smoke/Friends: a classical example for MLNs

This simple DeepDive example is based on a classical [Markov Logic Networks](#) example to show probabilistic inference and factor graphs functionalities of DeepDive. You can read more about [probabilistic inference and factor graphs](#) in our detailed documentation.

Inference rules

The objectives in this example are to infer whether a person smokes, and whether a person has cancer with a some probability using the factor $A \Rightarrow B$. Here, $A \Rightarrow B$ reads "A implies B", meaning that if A is true, then B is also true.

We introduce two rules:

1. If person A smokes, then A might have a cancer.
2. If two people A and B are friends and A smokes, then B might also smoke.

These rules can be written in DDlog as follows (in the [app.ddlog file](#)):

```
person (  
  person_id bigint,  
  name text  
).  
  
person_has_cancer? (  
  person_id bigint  
).  
  
person_smokes? (  
  person_id bigint  
).  
  
friends (  
  person_id1 bigint,  
  person_id2 bigint,  
  friend_since bigint  
).
```

```

    person_id bigint,
    friend_id bigint
).

@weight(0.5)
person_smokes(p) => person_has_cancer(p) :-
    person(p, _).

@weight(0.4)
person_smokes(p1) => person_smokes(p2) :-
    person(p1, _), person(p2, _), friends(p1, p2).

```

Setup

Before running the example, please check that DeepDive has been properly [installed](#) and the [necessary files](#) ([app.ddlog](#), [db.url](#), and [deepdive.conf](#)) and [directories](#) ([input/](#)) associated with this example are stored in the current working directory. Input directory should have [the data files](#) ([friends.tsv](#), [person_has_cancer.tsv](#), [person_smokes.tsv](#), and [person.tsv](#)). In order to use DeepDive, a database instance must be running to accept requests, and the database location must be specified in the [db.url](#). You can refer to the [tutorial](#) for further detail.

Running

Now you are ready to run the example. First, you have to compile the code using the following command.

```
deepdive compile
```

Once it has compiled with no error, you can run the following command to see the list of deepdive targets.

```
deepdive do
```

To run the entire pipeline you can run the following command.

```
deepdive run
```

This will display a plan for deepdive to run your pipeline. To start the pipeline, exit the editor with [:wq](#) command.

Results

Once the pipeline has completed running, you can view the results in the database using SQL or DDlog queries. The entire database should look like this:

List of relations				
Schema	Name			Type
pe	Owner	Size	Description	
-----+-----				
-+-----+-----+-----+-----				
public	dd_graph_variables_holdout			
table	user	0 bytes		
public	dd_graph_variables_observation			
table	user	0 bytes		
public	dd_graph_weights			
view	user	0 bytes		v
public	dd_inference_result_variables			
table	user	8192 bytes		
public	dd_factors_inf_imply_person_smokes_person_has_cancer			
table	user	8192 bytes		
public	dd_factors_inf_imply_person_smokes_person_smokes			
table	user	8192 bytes		
public	dd_weights_inf_imply_person_smokes_person_has_cancer			
table	user	16 kB		
public	dd_weights_inf_imply_person_smokes_person_smokes			
table	user	16 kB		
public	friends			
table	user	8192 bytes		ta
public	person			
table	user	16 kB		ta
public	person_has_cancer			
table	user	8192 bytes		t
public	person_has_cancer_calibration			
view	user	0 bytes		
public	person_has_cancer_inference			
view	user	0 bytes		

```

public | person_smokes | t
able | user | 8192 bytes |

public | person_smokes_calibration |
view | user | 0 bytes |

public | person_smokes_inference | v
iew | user | 0 bytes |

(16 rows)

```

Tables `person`, `friends`, `person_has_cancer`, and `person_smokes` hold the input data we prepared under the `input/` directory. To see what DeepDive inferred from our data, you can look at `person_smokes_inference` and `person_has_cancer_inference`. The two views should look like the following:

```

deepdive sql "SELECT * FROM person_smokes_inference"

```

person_id	dd_id	label	category	expectation
4	9		1	0.643
2	7		1	0.506
6	11		1	0.468
5	10		1	0.451

(4 rows)

```

deepdive sql "SELECT * FROM person_has_cancer_inference"

```

person_id	dd_id	label	category	expectation
3	2		1	0.635
1	0		1	0.614
6	5		1	0.57
2	1		1	0.563
4	3		1	0.563
5	4		1	0.551

(6 rows)

The `dd_id` column is for internal usage and can be ignored by the user and `person_id` is the user defined identifier in the input data. You can see that DeepDive uses the given data and inference

rules to predict the probability of the person being a smoker or having cancer in the **expectation** column.