

DeepDive Tutorial

Extracting mentions of spouses from the news

NOTE: we recommend using the [Jupyter Notebook version of this tutorial](#) that can be [easily launched inside Docker containers](#).

In this tutorial, we show an example of a prototypical task that DeepDive is often applied to: **extraction of structured information from unstructured or "dark" data** such as web pages, text documents, images, etc. While DeepDive can be used as a more general platform for statistical learning and data processing, most of the tooling described herein has been built for this type of use case, based on our experience of successfully applying DeepDive to [a variety of real-world problems of this type](#).

In this setting, our goal is to take in a set of unstructured (and/or structured) inputs, and populate a relational database table with extracted outputs, along with marginal probabilities for each extraction representing DeepDive's confidence in the extraction. More formally, we write a DeepDive application to extract mentions of *relations* and their constituent *entities* or *attributes*, according to a specified schema; this task is often referred to as **relation extraction**.^{*} Accordingly, we'll walk through an example scenario where we wish to extract mentions of two people being spouses from news articles.

The high-level steps we'll follow are:

1. **Data processing.** First, we'll load the raw corpus, add NLP markups, extract a set of *candidate* relation mentions, and a sparse *feature* representation of each.
2. **Distant supervision with data and rules.** Next, we'll use various strategies to provide *supervision* for our dataset, so that we can use machine learning to learn the weights of a model.
3. **Learning and inference: model specification.** Then, we'll specify the high-level configuration of our *model*.
4. **Error analysis and debugging.** Finally, we'll show how to use DeepDive's labeling, error analysis and debugging tools.

**Note the distinction between extraction of true, i.e., factual, relations and extraction of mentions of relations. In this tutorial, we do the latter, however DeepDive supports further downstream methods for tackling the former task in a principled manner.*

Whenever something isn't clear, you can always refer to [the complete example code at `examples/spouse/`](#) that contains everything shown in this document.

0. Preparation

First of all, make sure that [DeepDive has been installed](#).

Next, DeepDive will store all data—input, intermediate, output, etc.—in a relational database. Currently, Postgres, Greenplum, and MySQL are supported; however, Greenplum or Postgres are strongly recommended. To set the location of this database, we need to configure a URL in the `db.url` file, e.g.:

```
echo "postgresql://$USER@$HOSTNAME:5432/deepdive_spouse_$USER"  
>db.url
```

Note: DeepDive will drop and then create this database if run from scratch—beware of pointing to an existing populated one!

批注 [h1]: 使用
postgresql://localhost/deepdive_spouse_\$USER 不会报错

1. Data processing

In this section, we'll generate the traditional inputs of a statistical learning-type problem: candidate spouse relations, represented by a set of features, which we will aim to classify as *actual* relation mentions or not.

We'll do this in four basic steps:

1. Loading raw input data
2. Adding NLP markups
3. Extracting candidate relation mentions
4. Extracting features for each candidate

1.1. Loading raw input data

Our first task is to download and load the raw text of [a corpus of news articles provided by Signal Media](#) into an `articles` table in our database. We create a simple shell script that downloads and outputs the news articles in TSV format. DeepDive will automatically create the table, execute the script and load the table if we save it as:

```
input/articles.tsv.sh
```

The aforementioned script reads a sample of the corpus (provided as lines of JSON objects), and then using the `jq` language extracts the fields `id` (for document id) and `content` from each entry and converts those to TSV format.

Next, we need to declare the schema of this `articles` table in [our app.ddlog file](#); we add the following lines:

```
articles(  
  id      text,  
  content text
```

).

Then we compile our application, as we must do whenever we change `app.ddlog`:

```
deepdive compile
```

Finally, we tell DeepDive to execute the steps to load the `articles` table using the `input/articles.tsv.sh` script. You must have the [full corpus](#) downloaded.

```
deepdive do articles
```

Alternatively, a sample of 1000 articles can be loaded directly by issuing the following command: `bash deepdive load articles input/articles-1000.tsv.bz2`

DeepDive will output an execution plan, which will pop up in your default text editor. Save and exit to accept. DeepDive will run, creating the table and then fetching and loading the data.

After that finishes, we can take a look at the loaded data using the following `deepdive query` command, which enumerates the values for the `id` column of the `articles` table:

```
deepdive query '?- articles(id, _).'
```

```
id
```

```
-----  
340bb625-bb7e-49af-aa8d-781e5762f7a3  
fe4d045d-6c97-4ccf-a7d5-08aafb72a0db  
3585cd84-abd9-4f14-90ff-4fc23b45c84f  
eb571ae7-cd9f-4ad2-90a1-a74854574dc9  
6354c1a9-8f85-43c6-98d8-6b559d6727ef  
8994a0f1-7ca5-45a7-88c6-35bb98e89c95  
8b7fe0b3-ef78-4555-bb7a-5ad7007f9a93  
74ae2f47-ed24-4768-b8db-980f128061da  
89d7ee41-1b11-4f90-b1c9-e69c0145227d  
8f9d6a38-a828-40c5-b823-13447e96b842  
71984e17-d556-4828-93d3-18daf55a677f  
ce117e2b-73cb-48bc-9896-f9c9cb232a7e  
06b83cb4-1d03-42e9-8429-6cd272cd4a89  
f3397ce9-ba4f-4ddf-93c5-42333b98b907  
66ffaca4-4936-4844-a58e-b1a0b02c8df1  
84b24e59-a2ee-4eb4-a571-916d83ee475b
```

```
3fdcbae1-124e-46af-9773-2401683d19fd
e94a28ec-d335-46c5-aa7f-f14dfdc7258a
69a48149-d459-404f-adbe-6cd6049f71ee
ac5fbfd9-3f62-44bf-ab5a-36ffcd77adb1
70f5365b-51be-40f9-8fe4-380b6010e767
3f328bf8-45a3-441a-9426-2206a8b60b21
55021a16-6d82-4dcf-acef-21d1497910b1
fcd8d601-cce7-4735-92c0-eae086404d44
3a6277d7-8d1f-43bc-a69e-53f640a651a1
e1d7be87-2272-4091-89cb-173939b796a7
48eb9031-30b8-4689-987a-5ee3289448ee
84e7269f-3455-4af6-a1d4-541ae5b34e79
[...]
```

1.2. Adding NLP markups

Next, we'll use Stanford's [CoreNLP](#) natural language processing (NLP) system to add useful markups and structure to our input data. This step will split up our articles into sentences and their component *tokens* (roughly, the words). Additionally, we'll get *lemmas* (normalized word forms), *part-of-speech (POS) tags*, *named entity recognition (NER) tags*, and a dependency parse of the sentence. We declare the output schema of this step in `app.ddlog`:

```
sentences(
  doc_id      text,
  sentence_index int,
  sentence_text text,
  tokens      text[],
  lemmas      text[],
  pos_tags    text[],
  ner_tags    text[],
  doc_offsets int[],
  dep_types   text[],
  dep_tokens  int[]
```

```
).
```

Next we declare a DDlog function which takes in the `doc_id` and `content` for an article and returns rows conforming to the sentences schema we just declared, using the **user-defined function (UDF)** in `udf/nlp_markup.sh`. This UDF is a Bash script which calls [our own wrapper around CoreNLP](#). The CoreNLP library requires Java 8 to run.

```
function nlp_markup over (  
    doc_id text,  
    content text  
) returns rows like sentences  
implementation "udf/nlp_markup.sh" handles tsv lines.
```

Finally, we specify that this `nlp_markup` function should be run over each row from `articles`, and the output appended to `sentences`:

```
sentences += nlp_markup(doc_id, content) :-  
    articles(doc_id, content).
```

Again, to execute, we compile and then run:

```
deepdive compile  
deepdive do sentences
```

Now, if we take a look at a sample of the NLP markups, they will have tokens and NER tags that look like the following:

```
deepdive query '  
    doc_id, index, tokens, ner_tags | 5  
    ?- sentences(doc_id, index, text, tokens, lemmas, pos_tags,  
    ner_tags, _, _, _).
```

```
'  
  
doc_id | index |  
tok  
ens  
ner_tags  
-----+-----  
-----+-----  
-----  
-----+-----  
-----
```

Note that the previous steps—here, loading the articles—will *not* be re-run unless we specify that they should be, using, e.g.:

```
deepdive mark todo articles
```

1.3. Extracting candidate relation mentions

Mentions of people

Once again we first declare the schema:

```
person_mention(  
    mention_id    text,  
    mention_text  text,  
    doc_id        text,  
    sentence_index int,  
    begin_index   int,  
    end_index     int  
).
```

We will be storing each person as a row referencing a sentence with beginning and ending indexes. Again, we next declare a function that references a UDF and takes as input the sentence tokens and NER tags:

```
function map_person_mention over (  
    doc_id        text,  
    sentence_index int,  
    tokens        text[],  
    ner_tags      text[]  
    ) returns rows like person_mention  
    implementation "udf/map_person_mention.py" handles tsv lines.
```

We'll write a simple UDF in Python that will tag spans of contiguous tokens with the NER tag PERSON as person mentions (i.e., we'll essentially rely on CoreNLP's NER module). Note that we've already used a Bash script as a UDF, and indeed any programming language can be used. (DeepDive will just check the path specified in the top line, e.g., `#!/usr/bin/env python`.) However, DeepDive provides some convenient utilities for Python UDFs which handle all IO encoding/decoding. To write our UDF (`udf/map_person_mention.py`), we'll start by specifying that our UDF will handle TSV lines (as specified in the DDlog above). Additionally, we'll specify the exact type schema of both input and output, which DeepDive will check for us:

```
#!/usr/bin/env python
from deepdive import *

# for python 3 compatibility
try:
    xrange
except NameError:
    xrange = range

@tsj_extractor
@returns(lambda
    mention_id      = "text",
    mention_text     = "text",
    doc_id           = "text",
    sentence_index   = "int",
    begin_index      = "int",
    end_index        = "int",
    :[])
def extract(
    doc_id           = "text",
    sentence_index   = "int",
    tokens            = "text[]",
    ner_tags          = "text[]",
    ):
    """
    Finds phrases that are continuous words tagged with PERSON.
    """
    num_tokens = len(ner_tags)
    # find all first indexes of series of tokens tagged as PERSON
N
```



```

first_indexes = (i for i in xrange(num_tokens) if ner_tags
[i] == "PERSON" and (i == 0 or ner_tags[i-1] != "PERSON"))

for begin_index in first_indexes:

    # find the end of the PERSON phrase (consecutive tokens
tagged as PERSON)

    end_index = begin_index + 1

    while end_index < num_tokens and ner_tags[end_index] ==
"PERSON":

        end_index += 1

    end_index -= 1

    # generate a mention identifier

    mention_id = "%s_%d_%d_%d" % (doc_id, sentence_index, be
gin_index, end_index)

    mention_text = " ".join(map(lambda i: tokens[i], xrange
(begin_index, end_index + 1)))

    # Output a tuple for each PERSON phrase

    yield [
        mention_id,
        mention_text,
        doc_id,
        sentence_index,
        begin_index,
        end_index,
    ]

```

Above, we write a simple function which extracts and tags all subsequences of tokens having the NER tag "PERSON". Note that the `extract` function must be a generator (i.e., use a `yield` statement to return output rows).

Finally, we specify that the function will be applied to rows from the `sentences` table and append to the `person_mention` table:

```

person_mention += map_person_mention(
    doc_id, sentence_index, tokens, ner_tags
) :-

```

```
s, _, _, _).
    sentences(doc_id, sentence_index, _, tokens, _, _, ner_tag
```

Again, to run, just compile and execute as in previous steps:

```
deepdive compile && deepdive do person_mention
```

Now, the `person_mention` table should hold rows that look like the following:

```
deepdive query '
```

```
    name, doc, sentence, begin, end | 20
```

```
    ?- person_mention(p_id, name, doc, sentence, begin, end).
```

```
;
```

name			doc			sentenc		
e begin end								
-----+-----+-----								
-----+-----+-----								
Juliette Barnes			d9b82bc6-efa3-4c10-b595-8c51bbe27f3c					
1	5	6						
Hayden Panettiere			d9b82bc6-efa3-4c10-b595-8c51bbe27f3c					
1	8	9						
Shkreli			85fb02bf-6105-4dfb-995b-bd94a15a9e6d					
10	4	4						
Benjamin Davies			85fb02bf-6105-4dfb-995b-bd94a15a9e6d					
11	5	6						
Shkreli			85fb02bf-6105-4dfb-995b-bd94a15a9e6d					
12	9	9						
Alan Craze			ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6					
2	3	4						
Mary Shipstone			ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6					
4	0	1						
Maryam Alromisse			ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6					
4	6	7						
Yasser Alromisse			ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6					
5	3	4						
Mary			ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6					
6	9	9						
Lyndsey Shipstone			ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6					
6	15	16						

Alromisse	ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6	
7	7	7
Alromisse	ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6	
8	0	0
Craze	ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6	
10	1	1
Mary	ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6	
10	4	4
Alromisse	ddf3dfd2-cc21-46ca-a0e7-bd66eeb6bec6	
11	4	4
Chris Isaak	d9b2b908-190f-4010-a5b8-9a083f2b7fe5	
15	0	1
Guy Sebastian	d9b2b908-190f-4010-a5b8-9a083f2b7fe5	
15	16	17
Dannii Minogue	d9b2b908-190f-4010-a5b8-9a083f2b7fe5	
15	19	20
James Blunt	d9b2b908-190f-4010-a5b8-9a083f2b7fe5	
15	22	23
(20 rows)		

Mentions of spouses (pairs of people)

Next, we'll take all pairs of **non-overlapping person mentions that co-occur in a sentence with less than 5 people total**, and consider these as the set of potential ('candidate') spouse mentions. We thus filter out sentences with large numbers of people for the purposes of this tutorial; however, these could be included if desired. Again, to start, we declare the schema for our `spouse_candidate` table—here just the two names, and the two `person_mention` IDs referred to:

```
spouse_candidate(
  p1_id text,
  p1_name text,
  p2_id text,
  p2_name text
).
```

Next, for this operation we don't use any UDF script, instead relying entirely on **DDlog operations**. We simply construct a table of person counts, and then do a join with our filtering conditions. In DDlog this looks like:

```
num_people(doc_id, sentence_index, COUNT(p)) :-
    person_mention(p, _, doc_id, sentence_index, _, _).

spouse_candidate(p1, p1_name, p2, p2_name) :-
    num_people(same_doc, same_sentence, num_p),
    person_mention(p1, p1_name, same_doc, same_sentence, p1_begin, _),
    person_mention(p2, p2_name, same_doc, same_sentence, p2_begin, _),
    num_p < 5,
    p1_name != p2_name,
    p1_begin != p2_begin.
```

Again, to run, just compile and execute as in previous steps.

```
deepdive compile && deepdive do spouse_candidate
```

Now, the rows in the spouse candidates should look like the following:

```
deepdive query '
    name1, name2, doc, sentence | 20
?- spouse_candidate(p1, name1, p2, name2),
    person_mention(p1, _, doc, sentence, _, _).
'
```

name1	name2	doc	sentence
William Larnach	Eliza	ab9b3aa5-1b50-4dca-a179-fc2da32cd53e	40
Jonathan Kraft	Richard M. Berman	c94ecb16-3f32-42b3-8982-19c3ce8c0177	11
Jonathan Kraft	Brady	c94ecb16-3f32-42b3-8982-19c3ce8c0177	11

Jonathan Kraft	Roger Goodell	c94ecb16-3f32-42b3-8982-19c3ce8c0177	11
Nirmala Sitharaman	Biden Mumbai News.Net	e505b558-5db9-4d92-aec0-3b32341499b8	8
Nirmala Sitharaman	Joe Biden	e505b558-5db9-4d92-aec0-3b32341499b8	8
Nirmala Sitharaman	Husain Haqqani	e505b558-5db9-4d92-aec0-3b32341499b8	8
Cecil	Palmer	118e3e65-9264-4c3d-b9b1-c98a3aed3535	3
Mason	Mason	89ad2fcc-9cb5-41f1-9bd3-0123fcef3ac9	38
Palmer	Cecil	118e3e65-9264-4c3d-b9b1-c98a3aed3535	9
Cecil	Palmer	118e3e65-9264-4c3d-b9b1-c98a3aed3535	9
Barker	Constance	ab9b3aa5-1b50-4dca-a179-fc2da32cd53e	18
Barker	Jill Moon	ab9b3aa5-1b50-4dca-a179-fc2da32cd53e	18
Steph Laberis	Marty	8b39759d-835b-4b69-84a4-8359453be9cb	91
Steph Laberis	Marty	8b39759d-835b-4b69-84a4-8359453be9cb	91
Heyde	Theodor Eicke	1b48ada9-ae5e-4315-b032-95539bdb31d2	9
Rustam Kupaisinov	Krister Petersson	7d429da5-d18a-41ef-a416-9c15b088952e	2
Rustam Kupaisinov	Yury Zhukovsky	7d429da5-d18a-41ef-a416-9c15b088952e	2
Jimmy Tarbuck	Coleen Nolan	8bb7e4ae-828e-417a-8a15-b182a6bab1e1	4
Jimmy Tarbuck	Tarbuck	8bb7e4ae-828e-417a-8a15-b182a6bab1e1	4

(20 rows)

1.4. Extracting features for each candidate

Finally, we will extract a set of **features** for each candidate:

```
spouse_feature(  
    p1_id  text,  
    p2_id  text,  
    feature text  
).
```

The goal here is to represent each spouse candidate mention by a set of attributes or **features** which capture at least the key aspects of the mention, and then let a machine learning model learn how much each feature is correlated with our decision variable ('is this a spouse mention?'). For those who have worked with machine learning systems before, note that we are using a sparse storage representation- you could think of a spouse candidate (**p1_id**, **p2_id**) as being represented by a vector of length $L = \text{COUNT}(\text{DISTINCT feature})$, consisting of all zeros except for at the indexes specified by the rows with key (**p1_id**, **p2_id**).

DeepDive includes an [automatic feature generation library](#), **DDLlib**, which we will use here. Although many state-of-the-art [applications](#) have been built using purely DDLlib-generated features, others can be used and/or added as well. To use DDLlib, we create a list of **ddl.lib.Word** objects, two **ddl.lib.Span** objects, and then use the function **get_generic_features_relation**, as shown in the following Python code for **udf/extract_spouse_features.py**:

```
#!/usr/bin/env python  
  
from deepdive import *  
  
import ddlib  
  
@tsj_extractor  
@returns(lambda  
    p1_id  = "text",  
    p2_id  = "text",  
    feature = "text",  
    :[])  
  
def extract(  
    p1_id      = "text",  
    p2_id      = "text",
```

```

        p1_begin_index = "int",
        p1_end_index   = "int",
        p2_begin_index = "int",
        p2_end_index   = "int",
        doc_id         = "text",
        sent_index      = "int",
        tokens         = "text[]",
        lemmas         = "text[]",
        pos_tags       = "text[]",
        ner_tags       = "text[]",
        dep_types      = "text[]",
        dep_parents     = "int[]",
    ):
        """
        Uses DDLIB to generate features for the spouse relation.
        """
        # Create a DDLIB sentence object, which is just a List of DD
        LIB Word objects
        sent = []
        for i,t in enumerate(tokens):
            sent.append(ddlib.Word(
                begin_char_offset=None,
                end_char_offset=None,
                word=t,
                lemma=lemmas[i],
                pos=pos_tags[i],
                ner=ner_tags[i],
                dep_par=dep_parents[i] - 1, # Note that as stored f
                rom CoreNLP 0 is ROOT, but for DDLIB -1 is ROOT
                dep_label=dep_types[i]))

```

```
# Create DDLIB Spans for the two person mentions
```

```
p1_span = ddlib.Span(begin_word_id=p1_begin_index, length=(p1_end_index-p1_begin_index+1))
```

```
p2_span = ddlib.Span(begin_word_id=p2_begin_index, length=(p2_end_index-p2_begin_index+1))
```

```
# Generate the generic features using DDLIB
```

```
for feature in ddlib.get_generic_features_relation(sent, p1_span, p2_span):
```

```
    yield [p1_id, p2_id, feature]
```

Note that getting the input for this UDF requires joining the `person_mention` and `sentences` tables:

```
function extract_spouse_features over (
```

```
    p1_id          text,
```

```
    p2_id          text,
```

```
    p1_begin_index int,
```

```
    p1_end_index   int,
```

```
    p2_begin_index int,
```

```
    p2_end_index   int,
```

```
    doc_id         text,
```

```
    sent_index     int,
```

```
    tokens         text[],
```

```
    lemmas         text[],
```

```
    pos_tags       text[],
```

```
    ner_tags       text[],
```

```
    dep_types      text[],
```

```
    dep_tokens     int[]
```

```
) returns rows like spouse_feature
```

```
implementation "udf/extract_spouse_features.py" handles tsv lines.
```

```
spouse_feature += extract_spouse_features(
```



```

    p1_id, p2_id, p1_begin_index, p1_end_index, p2_begin_index,
    p2_end_index,

    doc_id, sent_index, tokens, lemmas, pos_tags, ner_tags, dep
_types, dep_tokens
) :-

    person_mention(p1_id, _, doc_id, sent_index, p1_begin_inde
x, p1_end_index),

    person_mention(p2_id, _, doc_id, sent_index, p2_begin_inde
x, p2_end_index),

    sentences(doc_id, sent_index, _, tokens, lemmas, pos_tags,
ner_tags, _, dep_types, dep_tokens).

```

Again, to run, just compile and execute as in previous steps.

```
deepdive compile && deepdive do spouse_feature
```

If we take a look at a sample of the extracted features, they will look roughly like the following:

```
deepdive query '| 20 ?- spouse_feature(_, _, feature).'
```

```

feature
-----
WORD_SEQ[accepted a plea deal and testified against]
LEMMA_SEQ[accept a plea deal and testify against]
NER_SEQ[0 0 0 0 0 0 0]
POS_SEQ[VBD DT NN NN CC VBD IN]
W_LEMMA_L_1_R_1[.][.]
W_NER_L_1_R_1[0][0]
W_LEMMA_L_2_R_1[Gissendaner .][.]
W_NER_L_2_R_1[PERSON 0][0]
W_LEMMA_L_3_R_1[against Gissendaner .][.]
W_NER_L_3_R_1[0 PERSON 0][0]
NGRAM_1[accept]
NGRAM_2[accept a]
NGRAM_3[accept a plea]
NGRAM_1[a]
NGRAM_2[a plea]

```

```
NGRAM_3_[a plea deal]
```

```
NGRAM_1_[plea]
```

```
NGRAM_2_[plea deal]
```

```
NGRAM_3_[plea deal and]
```

```
NGRAM_1_[deal]
```

```
(20 rows)
```

Now we have generated what looks more like the standard input to a machine learning problem—a set of objects, represented by sets of features, which we want to classify (here, as true or false mentions of a spousal relation). However, we **don't have any supervised labels** (i.e., a set of correct answers) for a machine learning algorithm to learn from! In most real world applications, a sufficiently large set of supervised labels is *not* available. With DeepDive, we take the approach sometimes referred to as *distant supervision or data programming*, where we instead generate a **noisy set of labels using a mix of mappings from secondary datasets and other heuristic rules**.

2. Distant supervision with data and rules

In this section, we'll use *distant supervision* (or '*data programming*') to provide a noisy set of labels for candidate relation mentions, with which we will train a machine learning model.

We'll describe two basic categories of approaches:

1. Mapping from secondary data for distant supervision
2. Using heuristic rules for distant supervision

Finally, we'll describe a simple majority-vote approach to resolving multiple labels per example, which can be implemented within DDlog.

2.1. Mapping from secondary data for distant supervision

First, we'll try using an external structured dataset of known married couples, from [DBpedia](#), to distantly supervise our dataset. We'll download the relevant data, and then map it to our candidate spouse relations.

Extracting and downloading the DBpedia data

Our goal is to first extract a collection of known married couples from DBpedia and then load this into the `spouses_dbpedia` table in our database. To extract known married couples, we use the DBpedia dump present in [Google's BigQuery platform](#). First we extract the URI, name and spouse information from the DBpedia `person` table records in BigQuery for which the field `name` is not NULL. We use the following query:

```
SELECT URI, name, spouse
FROM [fh-bigquery:dbpedia.person]
where name <> "NULL"
```

We store the result of the above query in a local project table `dbpedia.validnames` and perform a self-join to obtain the pairs of married couples.

```
SELECT t1.name, t2.name
FROM [dbpedia.validnames] AS t1
JOIN EACH [dbpedia.validnames] AS t2
ON t1.spouse = t2.URI
```

The output of the above query is stored in a new table named `dbpedia.spouseraw`. Finally, we use the following query to remove symmetric duplicates.

```
SELECT p1, p2
FROM (SELECT t1_name as p1, t2_name as p2 FROM [dbpedia.spouseraw]),
     (SELECT t2_name as p1, t1_name as p2 FROM [dbpedia.spouseraw])
WHERE p1 < p2
```

The output of this query is stored in a local file. The file contains duplicate rows (BigQuery does not support `distinct`). It also contains noisy rows where the name field contains a string where the given name family name and multiple aliases were concatenated and reported in a string including the characters `{` and `}`. Using the Unix commands `sed`, `sort` and `uniq` we first remove the lines containing characters `{` and `}` and then duplicate entries. This results in an input file `spouses_dbpedia.csv` containing 6,126 entries of married couples.

Loading DBpedia data to database

We [compress and store](#) `spouses_dbpedia.csv` under the path:

```
input/spouses_dbpedia.csv.bz2
```

Notice that for DeepDive to load the data to the corresponding database table, the name of the input data again has to be stored in the directory `input/` and has the same name as the target database table. To load the data we execute the command:

```
deepdive do spouses_dbpedia
```

Now the database should include tuples that look like the following:

```
deepdive query '| 20 ?- spouses_dbpedia(name1, name2).'
```

name1	name2
-----+-----	
A. A. Gill	Amber Rudd
Aamir Ali Malik	Sanjeeda Shaikh
Abimael Guzmán	Augusta la Torre
Abraham Jacobi	Mary Corinna Putnam Jacobi
Addison Adrienne Forbes Montgomery	Derek Christopher Shepherd
Alan Crosland	Elaine Hammerstein
Albert	Anna Marie of Brunswick-Lüneburg
Albert	Archduchess Margarethe Klementine of Austria
Albert	Dorothea
Albert	Isabella Clara Eugenia
Albert	Karoline Friederike Franziska Stephanie Amalie Cecilie
Albert	Richardis of Schwerin
Aleksander Ludwik Radziwiłł	Katarzyna Eugenia Tyszkiewicz
Alessia Merz	Fabio Bazzani
Alexis Denisof	Alyson Hannigan
Alfonso XII	Maria Christina of Austria
Alice of Namur	Baldwin IV Count of Hainaut
Amine Gemayel	Joyce Gemayel
Andrés Pastrana Arango	Nohra Puyana Bickenbach

批注 [h2]: 需要先在 app.ddlog 中申明 spouses_dbpedia, 否则会报错 Unknown target

Andrew Cymek | Brigitte Kingsley
(20 rows)

Supervising spouse candidates with DBpedia data

First we'll declare a new table where we'll store the labels (referring to the spouse candidate mentions), with an integer value (**True=1, False=-1**) and a description (**rule_id**):

```
spouse_label(  
  p1_id text,  
  p2_id text,  
  label int,  
  rule_id text  
).
```

Next we'll implement a simple distant supervision rule which labels any spouse mention candidate with a pair of names appearing in DBpedia as true:

```
# distant supervision using data from DBpedia  
spouse_label(p1,p2, 1, "from_dbpedia") :-  
  spouse_candidate(p1, p1_name, p2, p2_name), spouses_dbpedia  
(n1, n2),  
  [ lower(n1) = lower(p1_name), lower(n2) = lower(p2_name) ;  
    lower(n2) = lower(p1_name), lower(n1) = lower(p2_name) ].
```

It should be noted that there are many clear ways in which this rule could be improved (fuzzy matching, more restrictive conditions, etc.), but this serves as an example of one major type of distant supervision rule.

2.2. Using heuristic rules for distant supervision

We can also create a supervision rule which does not rely on any secondary structured dataset like DBpedia, but instead just uses some heuristic. We set up a DDlog function, **supervise**, which uses a UDF containing several heuristic rules over the mention and sentence attributes:

```
function supervise over (  
  p1_id text, p1_begin int, p1_end int,  
  p2_id text, p2_begin int, p2_end int,
```

```

        doc_id      text,
        sentence_index int,
        sentence_text text,
        tokens      text[],
        lemmas      text[],
        pos_tags    text[],
        ner_tags    text[],
        dep_types   text[],
        dep_tokens  int[]
    ) returns (
        p1_id text, p2_id text, label int, rule_id text
    )
    implementation "udf/supervise_spouse.py" handles tsv lines.

spouse_label += supervise(
    p1_id, p1_begin, p1_end,
    p2_id, p2_begin, p2_end,
    doc_id, sentence_index, sentence_text,
    tokens, lemmas, pos_tags, ner_tags, dep_types, dep_token_in
daxes
) :-
    spouse_candidate(p1_id, _, p2_id, _),
    person_mention(p1_id, p1_text, doc_id, sentence_index, p1_b
egin, p1_end),
    person_mention(p2_id, p2_text, _, _, p2_beg
in, p2_end),
    sentences(
        doc_id, sentence_index, sentence_text,
        tokens, lemmas, pos_tags, ner_tags, _, dep_types, dep_to
ken_indexes
    ).

```

The Python UDF named `udf/supervise_spouse.py` contains several heuristic rules:

- Candidates with person mentions that are too far apart in the sentence are marked as false.
- Candidates with person mentions that have another person in between are marked as false.
- Candidates with person mentions that have words like "wife" or "husband" in between are marked as true.
- Candidates with person mentions that have "and" in between and "married" after are marked as true.
- Candidates with person mentions that have familial relation words in between are marked as false.

```
#!/usr/bin/env python
from deepdive import *
import random
from collections import namedtuple

SpouseLabel = namedtuple('SpouseLabel', 'p1_id, p2_id, label, type')

@tsj_extractor
@returns(lambda
    p1_id = "text",
    p2_id = "text",
    label = "int",
    rule_id = "text",
    :[])

# heuristic rules for finding positive/negative examples of spouse relationship mentions
def supervise(
    p1_id="text", p1_begin="int", p1_end="int",
    p2_id="text", p2_begin="int", p2_end="int",
    doc_id="text", sentence_index="int",
```

```

        tokens="text[]", lemmas="text[]", pos_tags="text[]", ne
ner_tags="text[]",
        dep_types="text[]", dep_token_indexes="int[]",
    ):

    # Constants
    MARRIED = frozenset(["wife", "husband"])
    FAMILY = frozenset(["mother", "father", "sister", "brother",
                        "brother-in-law"])
    MAX_DIST = 10

    # Common data objects
    p1_end_idx = min(p1_end, p2_end)
    p2_start_idx = max(p1_begin, p2_begin)
    p2_end_idx = max(p1_end, p2_end)
    intermediate_lemmas = lemmas[p1_end_idx+1:p2_start_idx]
    intermediate_ner_tags = ner_tags[p1_end_idx+1:p2_start_idx]
    tail_lemmas = lemmas[p2_end_idx+1:]
    spouse = SpouseLabel(p1_id=p1_id, p2_id=p2_id, label=None,
                        type=None)

    # Rule: Candidates that are too far apart
    if len(intermediate_lemmas) > MAX_DIST:
        yield spouse._replace(label=-1, type='neg:far_apart')

    # Rule: Candidates that have a third person in between
    if 'PERSON' in intermediate_ner_tags:
        yield spouse._replace(label=-1, type='neg:third_person_
between')

    # Rule: Sentences that contain wife/husband in between

```



```

# (<P1>)([ A-Za-z]+)(wife|husband)([ A-Za-z]+)(<P2>)
if len(MARRIED.intersection(intermediate_lemmas)) > 0:
    yield spouse._replace(label=1, type='pos:wife_husband_b
etween')

# Rule: Sentences that contain and ... married
# (<P1>)(and)?(<P2>)([ A-Za-z]+)(married)
if ("and" in intermediate_lemmas) and ("married" in tail_lemmas):
    yield spouse._replace(label=1, type='pos:married_after
')

# Rule: Sentences that contain familial relations:
# (<P1>)([ A-Za-z]+)(brother|stster|father|mother)
([ A-Za-z]+)(<P2>)
if len(FAMILY.intersection(intermediate_lemmas)) > 0:
    yield spouse._replace(label=-1, type='neg:familial_betw
een')

```

Note that the rough theory behind this approach is that we don't need high-quality (e.g., hand-labeled) supervision to learn a high quality model. Instead, **using statistical learning, we can in fact recover high-quality models from a large set of low-quality or *noisy* labels.**

2.3. Resolving multiple labels per example with majority vote

Finally, we implement a very simple majority vote procedure, all in DDlog, for resolving scenarios where a single spouse candidate mention has multiple conflicting labels. First, we sum the labels (which are all -1, 0, or 1):

```

spouse_label_resolved(p1_id, p2_id, SUM(vote)) :- spouse_label
(p1_id, p2_id, vote, rule_id).

```

Then, we simply threshold and add these labels to our decision variable table `has_spouse` (see next section for details here):

```

has_spouse(p1_id, p2_id) = if 1 > 0 then TRUE
else if 1 < 0 then FALSE

```

```
else NULL end :- spouse_label_resolved(p1_id, p2_id, 1).
```

We additionally make sure that all spouse candidate mentions *not* labeled by a rule are also included in this table:

```
has_spouse(p1, p2) = NULL :- spouse_candidate(p1, _, p2, _).
```

Once again, to execute all of the above, just run the following command:

```
deepdive compile && deepdive do has_spouse
```

Recall that **deepdive do** will execute all upstream tasks as well, so this will execute all of the previous steps!

Now, we can take a brief look at how many candidates are supervised by different rules, which will look something like the table below. Obviously, the counts will vary depending on your input corpus.

```
deepdive query 'rule, @order_by COUNT(1) ?- spouse_label(p1,p2, label, rule).'
```

rule	COUNT(1)
-----+-----	
pos:married_after	556
from_dbpedia	1010
neg:familial_between	24944
pos:wife_husband_between	53032
neg:third_person_between	194898
neg:far_apart	305488
	731982
(7 rows)	

3. Learning and inference: model specification

Now, we need to specify the actual model that DeepDive will perform learning and inference over. At a high level, this boils down to specifying three things:

1. What are the *variables* of interest that we want DeepDive to predict for us?
2. What are the *features* for each of these variables?
3. What are the *connections* between the variables?

Once we have specified the model in this way, DeepDive will *learn* the parameters of the model (the weights of the features and potentially the connections between variables), and then perform *statistical inference* over the learned model to determine the probability that each variable of interest is true.

For more advanced users: we are specifying a *factor graph* where the features are unary factors, and then using SGD and Gibbs sampling for learning and inference. Further technical detail is available [here](#).

3.1. Specifying prediction variables

In our case, we have one variable to predict per spouse candidate mention, namely, **is this mention actually indicating a spousal relation or not?** In other words, we want DeepDive to predict the value of **a Boolean variable** for each spouse candidate mention, indicating whether it is true or not. We specify this in `app.ddlog` as follows:

```
has_spouse?(  
  p1_id text,  
  p2_id text  
).
```

DeepDive will predict not only the value of these variables, but also the marginal probabilities, i.e., the confidence level that DeepDive has for each individual prediction.

3.2. Specifying features

Next, we indicate (i) that each `has_spouse` variable will be connected to the features of the corresponding `spouse_candidate` row, (ii) that we wish DeepDive to learn the weights of these features from our distantly supervised data, and (iii) that the weight of a specific feature across all instances should be the same, as follows:

```
@weight(f)  
has_spouse(p1_id, p2_id) :-  
  spouse_candidate(p1_id, _, p2_id, _),  
  spouse_feature(p1_id, p2_id, f).
```

3.3. Specifying connections between variables

Finally, we can specify dependencies between the prediction variables, with either learned or given weights. Here, we'll specify two such rules, with fixed (given) weights that we specify. First, we define a *symmetry* connection, namely specifying that if the model thinks a person

mention **p1** and a person mention **p2** indicate a spousal relationship in a sentence, then it should also think that the reverse is true, i.e., that **p2** and **p1** indicate one too:

```
@weight(3.0)
has_spouse(p1_id, p2_id) => has_spouse(p2_id, p1_id) :-
    spouse_candidate(p1_id, _, p2_id, _).
```

Next, we specify a rule that the model should be strongly biased towards finding one marriage indication per person mention. We do this inversely, using a negative weight, as follows:

```
@weight(-1.0)
has_spouse(p1_id, p2_id) => has_spouse(p1_id, p3_id) :-
    spouse_candidate(p1_id, _, p2_id, _),
    spouse_candidate(p1_id, _, p3_id, _).
```

3.4. Performing learning and inference

Finally, to perform learning and inference using the specified model, we need to run the following command:

```
deepdive compile && deepdive do probabilities
```

This will ground the model based on the data in the database, learn the weights, infer the expectations or marginal probabilities of the variables in the model, and then load them back to the database.

Let's take a look at the probabilities inferred by DeepDive for the **has_spouse** variables.

```
deepdive sql "SELECT p1_id, p2_id, expectation FROM has_spouse_inference ORDER BY random() LIMIT 20"
```

p1_id	p2_id	expectation
1d1cff32-f332-41c3-9c18-57f44b5d7b02_4_12_12	1d1cff32-f332-41c3-9c18-57f44b5d7b02_4_6_7	0.171
5c467615-1dfc-4399-be91-6979cf6eade8_16_20_20	5c467615-1dfc-4399-be91-6979cf6eade8_16_22_22	0.952
bdecf4c6-81eb-4851-b30c-84f80ff58048_10_20_20	bdecf4c6-81eb-4851-b30c-84f80ff58048_10_26_26	0.001
fd2c5043-52ac-44e6-af50-9d3c32f6b4ce_49_24_25	fd2c5043-52ac-44e6-af50-9d3c32f6b4ce_49_29_30	0.034

f55c8585-893f-4e06-b904-9a1af91586c2_17_6_7	f55c8585-893f-4e06-b904-9a1af91586c2_17_2_3	0.107
b9e56701-b378-4049-af23-5473b4878174_8_39_40	b9e56701-b378-4049-af23-5473b4878174_8_13_13	0
17166b09-c02b-436e-82ec-bfd5afc5f23d_11_7_8	17166b09-c02b-436e-82ec-bfd5afc5f23d_11_5_5	0.163
15aa10bc-cd54-4956-9fb8-a61a850b86e3_50_17_18	15aa10bc-cd54-4956-9fb8-a61a850b86e3_50_9_15	0.993
c0e7a274-dac0-4913-96cc-2ec5209c52e0_4_1_1	c0e7a274-dac0-4913-96cc-2ec5209c52e0_4_40_40	0
15ea7ce9-b341-40b1-a0a9-b784ebd010fc_7_12_12	15ea7ce9-b341-40b1-a0a9-b784ebd010fc_7_5_5	0.04
f7ab24b8-ec6d-4293-9022-648818b7cbfe_17_9_9	f7ab24b8-ec6d-4293-9022-648818b7cbfe_17_1_1	0.017
b0430512-20d6-4c62-aa06-988a3bdaa1c3_30_37_37	b0430512-20d6-4c62-aa06-988a3bdaa1c3_30_27_27	0.018
147396b5-a728-4fc7-989b-5479c617d484_32_16_17	147396b5-a728-4fc7-989b-5479c617d484_32_5_6	0
72dff05a-4092-4724-b1c5-28e47d5626ce_17_5_5	72dff05a-4092-4724-b1c5-28e47d5626ce_17_3_3	0.992
c502650f-d0ca-4298-ab96-64ceebf70626_17_10_10	c502650f-d0ca-4298-ab96-64ceebf70626_17_8_8	0.533
a8391c1e-12d1-4178-9d79-cf69c94d58f8_13_1_1	a8391c1e-12d1-4178-9d79-cf69c94d58f8_13_5_5	0.038
7e5b8b2b-623d-423d-9da9-8ecb1f3a8c26_20_11_12	7e5b8b2b-623d-423d-9da9-8ecb1f3a8c26_20_0_1	0
954b6e8e-dce8-47e5-a769-c058ee6dbc00_2_15_15	954b6e8e-dce8-47e5-a769-c058ee6dbc00_2_12_12	0.059
71f4101b-42c5-4536-bee9-cffdd840da74_34_0_0	71f4101b-42c5-4536-bee9-cffdd840da74_34_2_2	0.978
09e8cfe7-06e9-4abd-8269-db8a90192552_50_0_0	09e8cfe7-06e9-4abd-8269-db8a90192552_50_10_11	0.002

(20 rows)

4. Error analysis and debugging

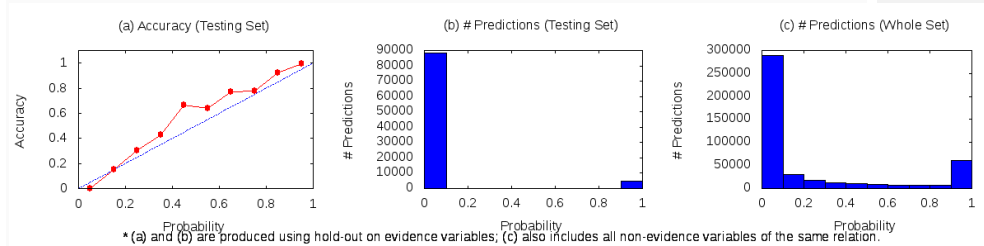
After finishing a pass of writing and running the DeepDive application, the first thing we want to see is how good the results are. In this section, we describe how **DeepDive's interactive tools** can be used for viewing the results as well as error analysis and debugging.

4.1. Calibration Plots

DeepDive provides *calibration plots* to see how well the expectations computed by the system are calibrated. The following command generates a plot for each variable under `run/model/calibration-plots/`.

```
deepdive do calibration-plots
```

It will produce a file `run/model/calibration-plots/has_spouse.png` that holds three plots as shown below:



Refer to the [full documentation on calibration data](#) for more detail on how to interpret the plots and take actions.

4.2. Browsing data with Mindbender

Mindbender is the name of the tool that provides an interactive user interface to DeepDive. It can be used for browsing any data that has been loaded into DeepDive and produced by it.

Browsing input corpus

We need to give hints to DeepDive about which part of the data we want to browse [using DDlog's annotation](#). For example, on the `articles` relation we declared earlier in `app.ddlog`, we can sprinkle some annotations such as `@source`, `@key`, and `@searchable`, as the following.

```
@source
articles(
  @key
  id text,
```

```
@searchable
```

```
content text
```

```
).
```

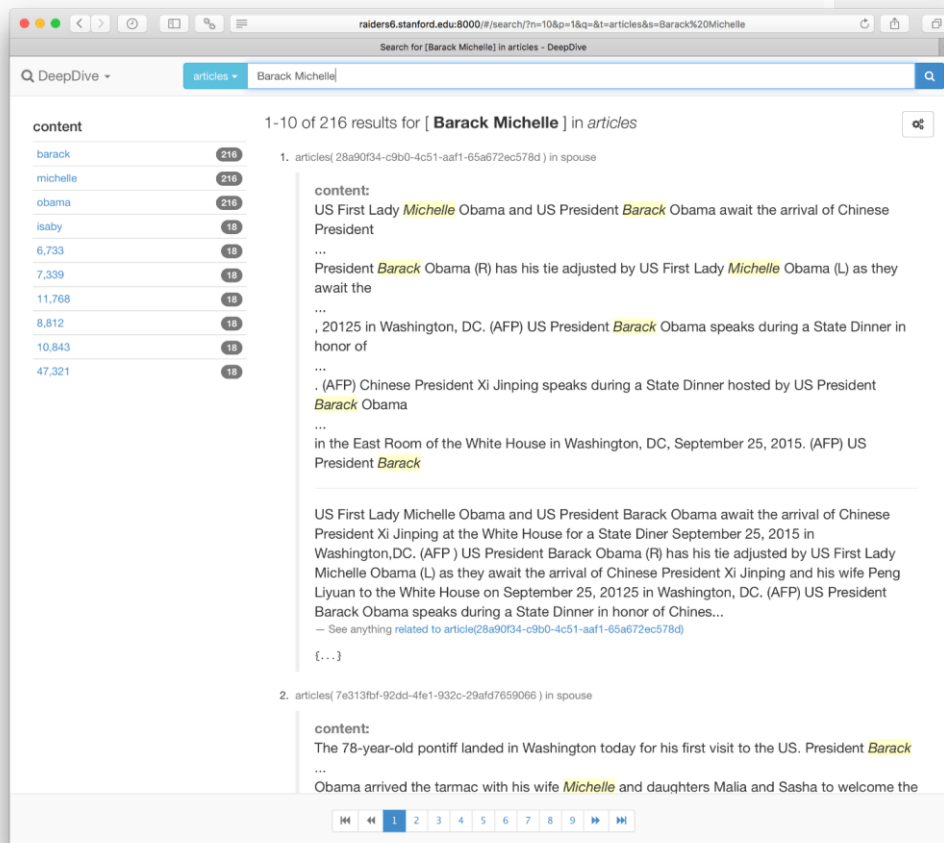
Next, if we run the following command, DeepDive will create and populate a search index according to these hints.

```
mindbender search update
```

To access the populated search index through a web browser, run:

```
mindbender search gui
```

Then, point your browser to the URL that appears after the command (typically <http://localhost:8000>) to see a view that looks like the following:



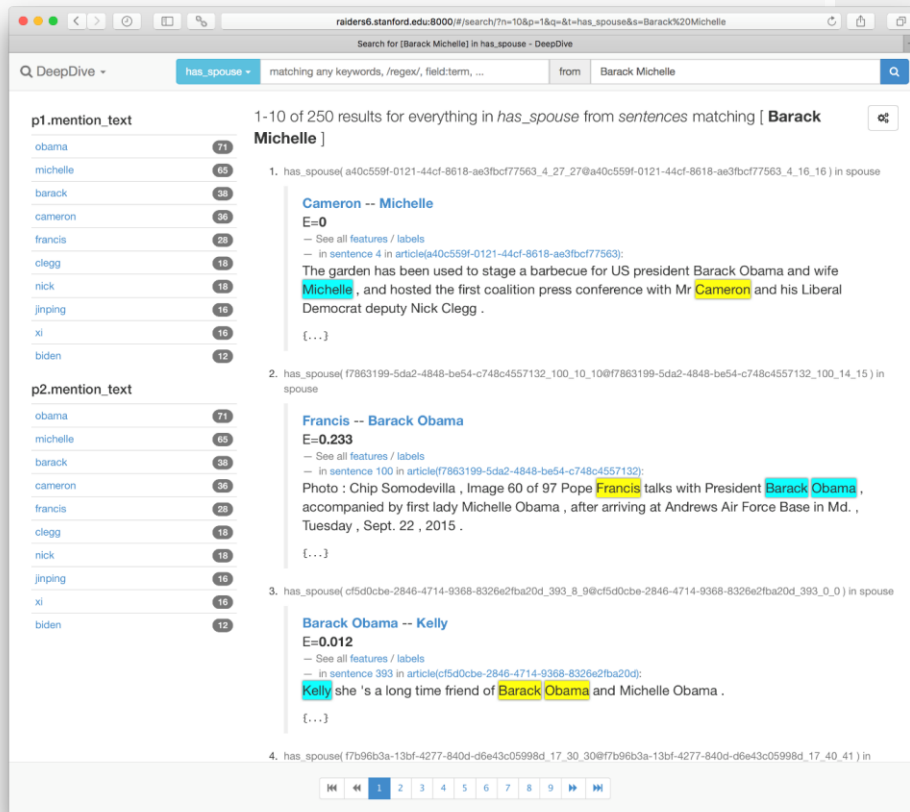
Browsing result data

To browse the results, we can add annotations to the inferred relations and how they relate to their source relations. For example, the `@extraction` and `@references` annotations in the following DDlog declaration tells DeepDive that the variable relation `has_spouse` is inferred from pairs of `person_mention`.

```
@extraction
has_spouse?(
  @key
  @references(relation="person_mention", column="mention_id",
    alias="p1")
  p1_id text,
  @key
  @references(relation="person_mention", column="mention_id",
    alias="p2")
  p2_id text
).
```

The relation `person_mention` as well as the relations it references should have similar annotations (see the [complete app.ddlog.code](#) for full detail).

Then, repeating the commands to update the search index and load the user interface will allow us to browse the expected marginal probabilities of `has_spouse` as well.



Customizing how data is presented

In fact, the screenshots above are showing the data presented using a [carefully prepared set of templates under mindbender/search-templates/](#). In these AngularJS templates, virtually anything you can program in HTML/CSS/JavaScript/CoffeeScript can be added to present the data that is ideal for human consumption (e.g., highlighted text spans rather than token indexes). Please see the [documentation about customizing the presentation](#) for further detail.

4.3. Estimating precision with Mindtagger

Mindtagger, which is part of the Mindbender tool suite, assists data labeling tasks to quickly assess the precision and/or recall of the extraction. We show how Mindtagger helps us perform a

labeling task to estimate the precision of the extraction. The necessary set of files shown below already exist [in the example under labeling/has_spouse-precision/](#).

Preparing a data labeling task

First, we can take a random sample of 100 examples from `has_spouse` relation whose expectation is higher than or equal to a 0.9 threshold as shown in [the following SQL query](#), and store them in [a file called `has_spouse.csv`](#).

```
deepdive sql eval "  
SELECT hsi.p1_id  
      , hsi.p2_id  
      , s.doc_id  
      , s.sentence_index  
      , hsi.label  
      , hsi.expectation  
      , s.tokens  
      , pm1.mention_text AS p1_text  
      , pm1.begin_index  AS p1_start  
      , pm1.end_index    AS p1_end  
      , pm2.mention_text AS p2_text  
      , pm2.begin_index  AS p2_start  
      , pm2.end_index    AS p2_end  
  
FROM has_spouse_inference hsi  
   , person_mention      pm1  
   , person_mention      pm2  
   , sentences            s  
  
WHERE hsi.p1_id          = pm1.mention_id  
      AND pm1.doc_id     = s.doc_id  
      AND pm1.sentence_index = s.sentence_index
```

```

AND hsi.p2_id      = pm2.mention_id
AND pm2.doc_id     = s.doc_id
AND pm2.sentence_index = s.sentence_index
AND      expectation >= 0.9

ORDER BY random()
LIMIT 100

" format=csv header=1 >labeling/has_spouse-precision/has_spouse.csv

```

We also prepare the `mindtagger.conf` and `template.html` files under `labeling/has_spouse-precision/` that look like the following:

```

title: Labeling task for estimating has_spouse precision
items: {
  file: has_spouse.csv
  key_columns: [p1_id, p2_id]
}
template: template.html

<mindtagger mode="precision">

  <template for="each-item">
    <strong title="item_id: "> -- </strong>
    with expectation <strong></strong> appeared in:
    <blockquote>
      <big mindtagger-word-array="item.tokens" array-format="
json">
        <mindtagger-highlight-words from="item.p1_start" to=
"item.p1_end" with-style="background-color: yellow;"/>
        <mindtagger-highlight-words from="item.p2_start" to=
"item.p2_end" with-style="background-color: cyan;"/>
      </big>
    </blockquote>
  </template>
</mindtagger>

```

```
</blockquote>
```

```
<div>
```

```
<div mindtagger-item-details></div>
```

```
</div>
```

```
</template>
```

```
<template for="tags">
```

```
<span mindtagger-adhoc-tags></span>
```

```
<span mindtagger-note-tags></span>
```

```
</template>
```

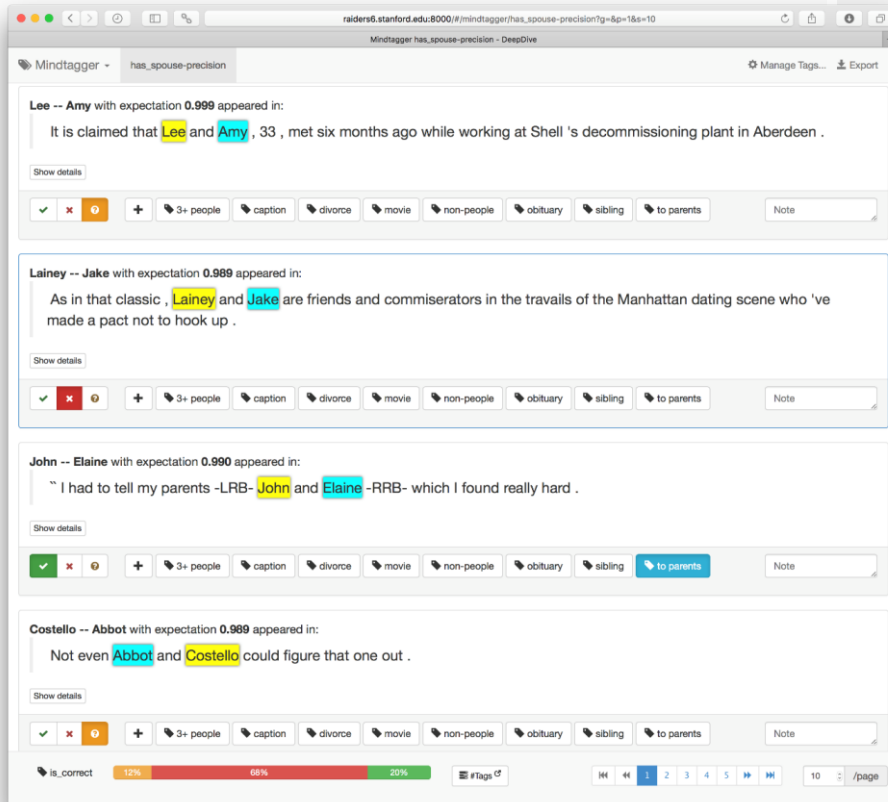
```
</mindtagger>
```

Labeling data with Mindtagger

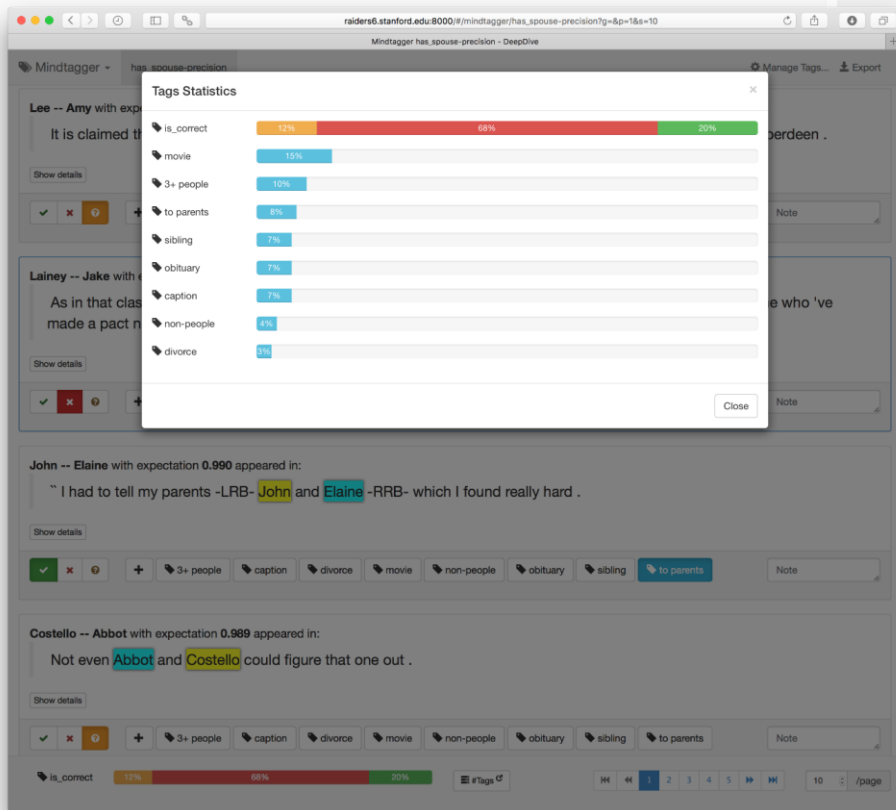
Mindtagger can then be started for the task using the following command:

```
mindbender tagger labeling/has_spouse-precision/mindtagger.conf
```

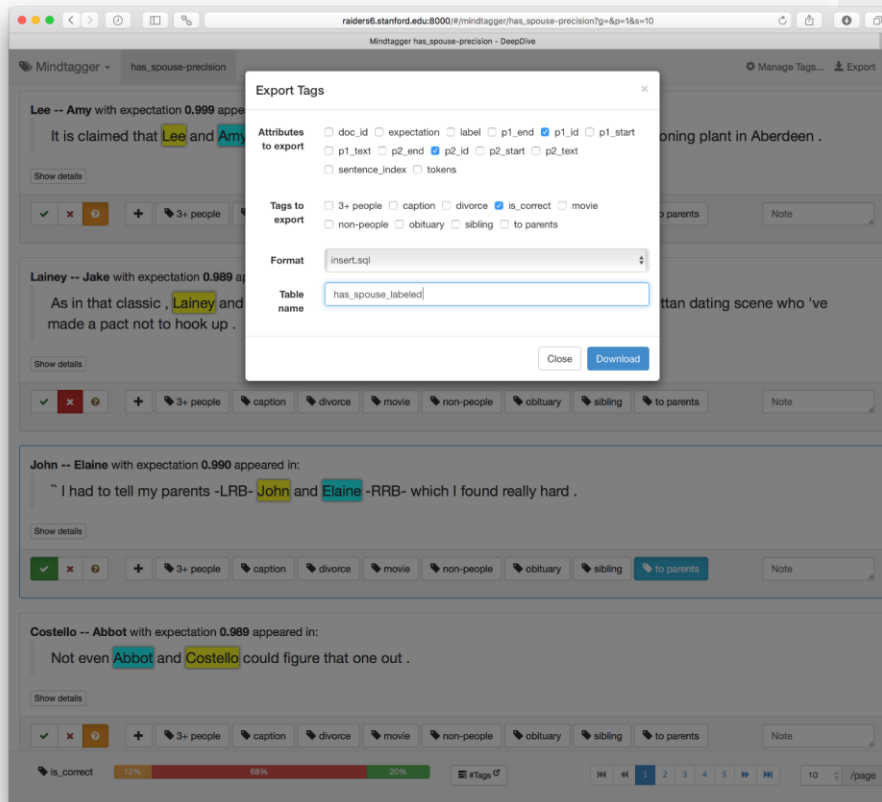
Then, point your browser to the URL that appears after the command (typically <http://localhost:8000>) to see a dedicated user interface for labeling data that looks like the following:



We can quickly label the sampled 100 examples using the intuitive user interface with buttons for correct/incorrect tags. It also supports keyboard shortcuts for entering labels and moving between items. (Press the **?** key to view all supported keys.) How many were labeled correct, as well as other tags, are shown in the "Tags" dropdown at the top right corner as shown below.



The collected tags can also be exported in various format for post-processing.



For further detail, see the [documentation about labeling data](#).

4.4. Monitoring statistics with Dashboard

Dashboard provides a way to monitor various descriptive statistics of the data products after each pass of DeepDive improvements. We can use a combination of SQL, any Bash script, and Markdown in each *report template* that produces a *report*, and we can produce a collection of them as a *snapshot* against the data extracted by DeepDive. Dashboard provides a structure to manage those templates and instantiate them in a sophisticated way using parameters. It provides a graphical interface for visualizing the collected statistics and trends as shown below. Refer to the [full documentation on Dashboard](#) to set up your own set of reports.

